

OpenCPI RCC Development Guide

DRAFT/Incomplete for review

Revision History

Revision	Description of Change	Date
0.1	Initial, from earlier doc sources	2015-10-31
0.9	Add all C++ content slave/proxy content, references to CDG	2016-01-27

Table of Contents

1 Introduction.....	5
1.1 References.....	6
2 Overview.....	7
2.1 Descriptive XML for RCC Workers.....	8
3 XML description files (OWD) for RCC workers.....	9
3.1 Attributes of a top-level RCCWorker element.....	10
3.1.1 Name attribute.....	10
3.1.2 Language attribute.....	10
3.1.3 ControlOperations attribute.....	10
3.1.4 Slave attribute.....	10
3.1.5 ExternMethods attribute (C language).....	10
3.2 Attributes of Port Elements in the OWD.....	11
3.2.1 Name attribute.....	11
3.2.2 MinBufferCount attribute.....	11
4 The RCC Worker Interface.....	12
4.1 Worker methods: called by the container, implemented by the worker.....	13
4.1.1 RCCWorker Structure Type – C language only.....	13
4.1.2 RCCResult Enumeration Type (C and C++).....	14
4.1.3 initialize — worker method.....	15
4.1.4 start — worker method.....	16
4.1.5 stop — worker method.....	16
4.1.6 release — worker method.....	17
4.1.7 afterConfigure — worker method.....	17
4.1.8 beforeQuery — worker method.....	18
4.1.9 run — worker method.....	18
4.2 Container methods, called by the worker.....	21
4.2.1 time/getTime — container method.....	21
4.2.2 setError — container method.....	21
4.2.3 getRunCondition — container method — C++ Language only.....	22
4.2.4 setRunCondition — container method — C++ Language only.....	22
4.2.5 The RCC::RunCondition C++ class.....	22
4.3 Port Management Data Members and Methods.....	25
4.3.1 advance — port method (C++), container method (C).....	25
4.3.2 hasBuffer — query a port for whether it has a current buffer-to-release.....	26
4.3.3 isConnected — query a port for being connected.....	26
4.3.4 request — port method (C++), container method (C).....	26
4.3.5 send — send an input buffer on an output port.....	27
4.3.6 setDefaultLength — set the default message length at an output port.....	27
4.3.7 setDefaultOpCode — set the default opcode at an output port.....	28
4.3.8 take — port method (C++), container method (C).....	28
4.4 Buffer Management Data Members and Methods.....	30
4.4.1 checkLength — check size of the buffer.....	30
4.4.2 data — access the raw contents of the buffer.....	30
4.4.3 length — retrieve the length of the message in a buffer – C++ only.....	31

4.4.4	maxLength — retrieve the maximum available space in the buffer.....	31
4.4.5	opCode — retrieve the opCode of the message in a buffer — C++ only.....	31
4.4.6	release — release a buffer.....	31
4.4.7	setInfo — set the metadata associated with a buffer — C++ only.....	32
4.4.8	setLength — set the length of the message in a buffer — C++ only.....	32
4.4.9	setOpCode — set the opCode for the message in a buffer — C++ only.....	32
4.4.10	topLength — retrieve the size of the single sequence in a message — C++ only.....	33
4.5	Accessing the Contents of Messages.....	34
4.6	Property Access Notifications.....	35
4.7	Controlling Slave Workers from Proxies (C++ only).....	37
4.8	Worker Dispatch Structures — C Language only.....	38
4.8.1	RCCDispatch Structure Type.....	38
5	Code Generation for RCC Workers.....	40
5.1	Namespace Management.....	41
5.2	Generated Data Types.....	42
5.2.1	The enumeration constants for the worker’s ports.....	42
5.2.2	The properties structure type.....	42
5.2.3	Structures for Message Payloads (C only).....	43
5.2.4	Worker Base Class (C++ only).....	44
6	RCC Local Services.....	45
6.1	RCC Local Services AEP as a small subset of POSIX and ISO-C.....	46
7	Summary of OpenCPI RCC authoring model.....	48
7.1	Worker code example.....	49
8	Glossary.....	55
9	References.....	56
10	List of Abbreviations and Acronyms.....	57

1 Introduction

This document specifies the OpenCPI RCC authoring model and describes how to write RCC workers, in C or C++. The term “RCC” is defined as “Resource-Constrained C Language”. This model is based on the C language and makes most design choices to minimize resources and thus to be appropriate for resource-constrained embedded systems. DSP processors with on-chip memories, micro-controllers, and some multi-core processors are natural targets for this authoring model. Of course the RCC model is a perfectly appropriate model for any general-purpose processor with a C compiler, when the developer is comfortable with the constraints of the C language.

This document also describes a C++ variant of this authoring model that takes advantage of the expressive power of the C++ language (which usually results in fewer SLOC – Source Lines of Code), at a modest cost of increased memory footprint. Both the C and C++ language variants are considered to be based on this one (RCC) authoring model since so many concepts and details are common.

This specification is based on the authoring model concept as defined in the OpenCPI Component Development Guide. That document introduces key concepts for all authoring models, including the configuration and lifecycle model of components, and the software execution model for most authoring models targeting general purpose software platforms (rather than FPGA or similar).

That document is a prerequisite for this one.

As with all authoring models, this model is intended to coexist and interoperate with the other existing models that are more appropriate for other processing technologies such as GPUs (the OCL authoring model), and FPGAs (the HDL authoring model). Other models and any unique aspects to their associated development workflow, are described in their own documents.

1.1 References

This document depends on several others. Primarily, it depends on the “OpenCPI Component Development Guide”, which describes concepts and definitions common to all OpenCPI authoring models. Since this RCC authoring model is based on the C language (specifically C90: ISO/IEC 9899:1990), it also depends on the ISO-C language reference manual and associated libraries. The exceptions to the C90 basis are the use of <stdint.h> from C99. The C++ authoring model is based on the language as defined in XXX, prior to Cxx11.

Table 1 - Table of Reference Documents

Title	Published By	Link
OpenCPI Component Development Guide	OpenCPI	githubURL
ISO C Language Specification	C Language	Public URL: http://www.iso.org/c90
ISO C++ Language Specification	C++	Pre cxx11 C++

2 Overview

RCC workers are hosted in a *container*, which is responsible for (1) loading, executing, controlling, and configuring the worker, (2) effecting data movement to and from the data ports of the worker, and (3) providing interfaces for the local services available to RCC workers.

When RCC workers are collocated, the container can make use of local zero-copy approaches to move data between them. For connections between workers in *different* containers, the containers move data between each other using some common data transport mechanism. Containers make use of a default data transport between the two devices unless explicitly configured to do otherwise.

The OpenCPI Component Development Guide contains sections for the general introduction to the control plane functionality of workers and containers, followed by the general execution model of software-based workers. The specifics of the RCC authoring model are included here, consisting of:

- ***container-to-worker interfaces:*** how the container calls the worker's entry points
- ***worker-to-container interfaces:*** how the worker calls the container's entry points
- ***the local services:*** how the worker uses local services and which ones are available.

2.1 Descriptive XML for RCC Workers.

As with other authoring models, the act of creating a component implementation (a.k.a. authoring a worker), includes writing source code as well as specifying certain characteristics of the implementation in a separate XML file. This XML file is called the OWD (for **O**penCPI **W**orker **D**escription). It describes any non-default constraints or behavior of this particular component implementation, and has attributes and information that are specific to the authoring model. The two items that all OWDs must specify is the component specification (OCS) being implemented, and the authoring model being used. The OWD aspects that are common to all authoring models are described in the OpenCPI Component Development Guide.

The model name for the RCC authoring model is “RCC”. The languages supported are C and C++. The top level XML element for the OWD for RCC workers is **RCCWorker**. The details of the OWD for RCC workers is described below in section 4.

3 XML description files (OWD) for RCC workers.

This section describes the format and structure of the RCC OWD: the XML document that refers to or includes a Component Specification (OCS), and which specifies the implementation-specific aspects of the worker. The OpenCPI Component Development Guide defines aspects common to all OWDs. This section defines the aspects of the RCC OWD – specific to RCC workers. The top level XML element for a RCC worker is `RCCWorker`. If all defaults are used, the entire OWD for an RCC worker can be:

```
<RCCWorker spec='myspec' />
```

The `RCCWorker` XML element contains information provided by someone creating an RCC worker based on a component specification (OCS). This OWD includes or references an OCS, and then describes implementation information about this particular RCC implementation of that OCS. Thus the `RCCWorker` element must either include as a child element a complete `ComponentSpec`, or include one by reference, using the “spec” attribute of the top-level element. For example, the “vsadd” implementation of the “vsadd-spec” OCS would reference the component specification this way:

```
<RCCWorker spec="vsadd-spec"  
  ---other attributes---  
>  
  ---other child elements---  
</RCCWorker>
```

The `RCCWorker` follows the specification of OWDs in general as specified elsewhere. This section only defines the aspects of the RCC OWD (`RCCWorker`) that are *not* common to all OWDs.

3.1 *Attributes of a top-level RCCWorker element*

3.1.1 *Name attribute*

The “Name” attribute of the component implementation is used to generate the programming language name of the worker (e.g. used in the name of and the contents of the RCC-specific generated header file). When not specified (the common case) the implementation name defaults to the name of the OWD file itself, without directories or extensions.

3.1.2 *Language attribute*

The “Language” attribute of the component implementation for RCC workers should have the value “c” or “c++”. The default is “c”.

3.1.3 *ControlOperations attribute*

If not specified, no control operation methods are implemented by the worker (only the `run` method). This information is used in the generation of the implementation skeleton to provide stubs for the specified control operation implementations. It is also stored in the runtime metadata associated with the worker, which allows the runtime system to avoid invoking control operations that have no implementation. This attribute is common to all OWDs.

3.1.4 *Slave attribute*

This attribute indicates that this worker is a proxy for another worker, and that other worker is named as the string value of this attribute. This attribute is only supported for C++. The name of the slave worker may have a package prefix (with periods) if the slave worker is not in the same namespace as this worker. The slave worker name must include the authoring model suffix. An example is:

```
<RCCWorker language='c++' slave='ocpi.devices.xyz_adc.hdl' />
```

This would indicate that this worker is a proxy for the HDL worker named `xyz_adc`, in the `ocpi.devices` package name scope.

3.1.5 *ExternMethods attribute (C language)*

This String attribute is used (when set to true), in the cases where the methods of the worker (whose addresses are placed in the `RCCDispatch` structure) should have external linkage scope rather than be local to the file (i.e. declared “static”). In this case, the string value of this attribute is used as the prefix to those external function names, rather than just being file-scoped static functions whose names are simply the method names (e.g. “initialize”, “start”, “run”, etc.).

3.2 Attributes of Port Elements in the OWD

The `port` child element of the `RCCWorker` specifies information about a data port specified in the OCS. It references an OCS `port` (or `dataInterfaceSpec`) element by its `Name` attribute. Thus the `Name` attribute of the `Port` element must match the `Name` attribute of a `Port` or `DataInterfaceSpec` element of the `ComponentSpec`. The `Port` element adds implementation-specific information about the port initially defined in that `ComponentSpec`.

A number of attributes available for the `port` element are common to all authoring models and are described in the Component Development Guide. These are usually used to override attributes inferred from the protocol associated with the data port.

3.2.1 Name attribute

This attribute specifies the name used to reference the `Port` or `DataInterfaceSpec` element in the `ComponentSpec` (OCS).

3.2.2 MinBufferCount attribute

This numeric attribute specifies the minimum number of message buffers required by the worker for a port. The Worker Interface allows the worker code (typically in the `run` method) to “take” a buffer from a port, and ask for a new buffer for that port while retaining ownership of the previous buffer from that port. This behavior requires that the infrastructure provide at least 2 buffers for that port.

Thus this attribute informs the infrastructure as to the minimum buffering requirements of the worker implementation for that port. The default is 1. *This attribute should not be used to “tune” the buffer count for performance but only specify the actual minimum requirements for the correct functioning of the worker.*

4 The RCC Worker Interface

This section defines the interface between the worker and its container: the API for this authoring model. This document uses the term “worker method” as shorthand and language-neutral term for what is a “member function” in C++ or a worker entry point function in C. The methods in C++ have an implicit “**this**” argument that is hidden by the language. In RCC C language workers, the first argument to all worker methods is an explicit “**self**” argument that is a pointer to a structure that contains context and state information for the worker. When discussing a worker's runtime behavior, the term “*worker*” is sometimes used as a runtime *instance* of the worker rather than the source code that is written for the worker as a component implementation.

RCC workers must avoid the prefixes “OCPI” and “RCC” for compile-time constants and types since they are used automatically by this authoring model. All defined data types for the RCC authoring model use the common prefix “RCC”. All macros are upper case, and all data types are capitalized and mixed (camel) case.

The worker interface (its methods) consists of control operation methods whose behavior is defined in the “Control Plane Introduction” section of the OpenCPI Component Development Guide, plus an additional **run** method that supports the event-driven execution model defined in the “Software Execution Model” of that same document. The **run** method is the only required method. All the other methods are optional. All processing of the worker occurs in the context of these methods.

The interface uses several basic integer types consistent with their CORBA C/C++ language mapping, to provide some compiler independence. The integer types are defined using the ISO C 99 `<stdint.h>` types. These basic types are: **uint8_t**, **uint16_t** and **uint32_t**. The type **RCCBoolean** is aliased to **uint8_t**, to be consistent with the defined size in property space and message content layouts. The **RCCOrdinal** type is an alias for **uint16_t**, and is used when ordinals are required (ports, operations, exceptions, properties).

There are two categories of methods:

- **Worker methods** that represent functionality of the worker, to be called by the container, and which may have default implementations. These include the “run” method, and the lifecycle control operation methods.
- **Container methods** that represent functionality of the container, to be called by the worker, such as changing run conditions, and accessing ports and buffers. All C language container methods are dispatched through function pointers in the **container** member of the **RCCWorker** structure.

Worker methods are described first, followed by container methods. Various data types used in worker and container methods are described in their respective sections. The concept of a “default” method for C-language workers is indicated by NULL pointers for those methods in the **RCCDispatch** structure. For C++, then are simply the methods in the base class.

4.1 Worker methods: called by the container, implemented by the worker.

This section describes the methods that a worker may implement. One is mandatory: **run**. All others are optional, and have default behavior. Other than “**run**”, the other methods are *control operations* that perform lifecycle state transitions described in the control plane introduction section of the OpenCPI Component Development Guide. All processing of the worker occurs in the context of these worker methods.

For C++, the code generation tools create a custom base class that the actual worker class inherits. Thus the derived class that implements the worker is declared by the worker author directly in their source code. This derived class inherits the custom base class, but can otherwise contain any other member functions and data members as long as they do not shadow certain members in the generated custom base class (which are specifically mentioned below).

4.1.1 RCCWorker Structure Type – C language only

This structure type (typedef name) represents the visible state of a worker. The container defines this structure (in **RCC_Worker.h**) with any content or member ordering as long as the documented members are supported. The structure members are written by either the container or the worker, but not both. Members written by the container are declared “const” to enhance error checking when compiling worker code. A pointer to this structure is the first argument to all C language worker methods.

Table 6 Members of the RCCWorker Structure for C Workers

Member Name	Member Data Type	Written by	Member Description
properties	void *const	container	A const pointer to the properties structure for the worker, whose layout is implied by the properties declared in the OCS and OWD. The value is NULL if there are no such properties.
memories	void *const *const	container	An array of const pointers to the memory resources requested by the worker in the memSizes member of the RCCDispatch structure. Any memories that are not read-only are initialized to zero before a worker executes any method.
memory	void *const	container	A pointer to the memory resource as requested in the memSize member of the RCCDispatch structure.
container	const RCCContainer	container	A dispatch table of container method/function pointers.
runCondition	RCCRunCondition *	worker	Initialized from the RCCDispatch runCondition member. Checked by container after calling the start method.
connectedPorts	RCCPortMask	container	A mask indicating which ports are connected. A worker can check this to see if an optional port is connected.
ports	RCCPort[]	varies by member	An array of RCCPort structures, indexed by port ordinals.

4.1.2 RCCResult Enumeration Type (C and C++)

The **RCCResult** type is an enumeration type used as the return value for all worker methods. It indicates to the container what to do when the worker method returns, as described in the following table:

Table 6: RCCResult Return Values

Enumeration Identifier	Description
RCC_OK	The worker method succeeded without error.
RCC_ERROR	The worker method did not succeed, but the error is not fatal to worker or container, thus the method may be retried if defined to allow this.
RCC_FATAL	The worker should never run again. It is non-functional. The container or other workers may be damaged. The worker is in an <i>unusable</i> state. The container may know that it, or other workers are protected from damage, but the worker indicates this condition in case there is no such protection.
RCC_DONE	The worker needs no more execution; a normal completion. The worker is entering the <i>finished</i> state.
RCC_ADVANCE	The worker is requesting that all ports that were ready when the run method was entered be <i>advanced</i> (applies only to the run method).
RCC_ADVANCE_DONE	The worker is requesting that all ports be advanced (run method only) and declaring that it is also “done”. The worker is entering the <i>finished</i> state.

These return values apply to each method as defined in their specific behavior. Some values are not valid results for all methods. When the result is **RCC_ERROR** or **RCC_FATAL**, the worker may have also set a descriptive error message via the **setError** container method described below.

4.1.3 *initialize* — worker method

This worker method implements the *initialize* control operation as defined in [CDG]. It is optional. A worker author should implement this method under any of the following conditions:

- There are initialization errors that can be reported to the container.
- There is significant processing work involved in the initialization.
- There are significant resource allocations performed during initialization.

If none of the above are true, this method can be unimplemented (and thus the default implementation will be used). Note that initial property settings are *not* available when the *initialize* method is called.

Using the `initialize` method will allow the `run` or `start` methods to be more consistent in execution time, since without an `initialize`, all one-time initialization must be done on the first invocation of `start`, or if `start` is unimplemented, the first invocation of `run`.

4.1.3.1 Synopsis

```
static RCCResult initialize(RCCWorker *self); // C language
RCCResult initialize(); // C++ language
```

4.1.3.2 Returns

This method shall return a `RCCResult` value.

If the initialization cannot succeed, it shall return `RCC_ERROR`. If the worker detects an error that would disable the implementation or its environment, it shall return `RCC_FATAL`. Otherwise it shall return `RCC_OK` if normal worker execution should proceed.

4.1.4 **start** — worker method

This method implements the **start** control operation as defined in [CDG]. It is optional. Upon successful completion, the worker is in the **operating** state or the **finished** state. This operation is called for the first time after initialization, and also may be called after a successful **stop** operation, to *resume* execution.

4.1.4.1 Synopsis

```
static RCCResult start(RCCWorker *self); // C language
RCCResult start(); // C++ language
```

4.1.4.2 Returns

This method shall return a `RCCResult` value.

If the **start** method cannot succeed, it shall return `RCC_ERROR`. If the worker detects an error that would disable the implementation or its environment, it shall return `RCC_FATAL`. It shall return `RCC_DONE` if no worker execution should proceed, to indicate that it is entering the **finished** state. Otherwise it shall return `RCC_OK` if normal worker execution should proceed.

Returning `RCC_DONE` indicates to the container that the worker will never require further execution, and provides this advice to the container to allow the container to take advantage of this fact and possibly release resources (such as I/O buffers) prior to the worker instance being destroyed or reused.

4.1.5 **stop** — worker method

This method implements the **stop** control operation as described in [CDG]. Upon successful completion, the worker is in the **suspended** state. It will be also called before the **release** method, and before destruction, if the worker is in the **operating** state.

4.1.5.1 Synopsis

```
static RCCResult stop(RCCWorker *self); // C language
RCCResult stop(); // C++ language
```

4.1.5.2 Returns

This method shall return a `RCCResult` value.

If the **stop** method cannot succeed, it shall return `RCC_ERROR`. If the worker detects an error that would disable the implementation or its environment, it shall return `RCC_FATAL`. Otherwise it shall return `RCC_OK`.

4.1.6 **release** — worker method

This method implements the *release* control operation as described in [CDG]. After successful completion, the worker instance is in the **exists** state.

4.1.6.1 Synopsis

```
static RCCResult release(RCCWorker *self); // C language
RCCResult release(); // C++ language
```

4.1.6.2 Returns

This method shall return a `RCCResult` value.

If the *release* method cannot succeed, it shall return `RCC_ERROR`. If the worker detects an error that would disable the implementation or its environment, or it is unable to return resources it allocated, it shall return `RCC_FATAL`. Otherwise it shall return `RCC_OK`.

4.1.7 **afterConfigure** — worker method

This optional method implements the *afterConfigure* control operation. It is called by the container after properties have been changed that have been specified in the OWD as requiring synchronization with the worker (using the `writeSync` attribute).

4.1.7.1 Synopsis

```
static RCCResult afterConfigure(RCCWorker* self); // C language
RCCResult afterConfigure(); // C++ language
```

4.1.7.2 Returns

This method shall return a `RCCResult` value.

If the *afterConfigure* method cannot succeed, it shall return `RCC_ERROR`. If the worker detects an error that would disable the implementation or its environment, it shall return `RCC_FATAL`. Otherwise it shall return `RCC_OK`.

4.1.8 *beforeQuery* — worker method

This optional method implements the *beforeQuery* control operation. It is called by the container before properties have been accessed that have been specified in the OWD as requiring synchronization with the worker (using the `readSync` attribute).

4.1.8.1 Synopsis

```
static RCCResult beforeQuery(RCCWorker *self); // C language
RCCResult beforeQuery(); // C++ language
```

4.1.8.2 Returns

This method shall return a `RCCResult` value.

If the *beforeQuery* method cannot succeed, it shall return `RCC_ERROR`. If the worker detects an error that would disable the implementation or its environment, it shall return `RCC_FATAL`. Otherwise it shall return `RCC_OK`.

4.1.9 *run* — worker method

The *run* method requests that the worker perform its normal computation. The container only calls this method when the worker's *run condition* is satisfied.

4.1.9.1 Synopsis

```
static RCCResult run(RCCWorker *self,
                    RCCBoolean timedOut,
                    RCCBoolean *newRunCondition); // C language
RCCResult run(bool timedOut); // C++ language
```

4.1.9.2 Behavior

The *run* method shall perform the worker's computational function and return a result. The *run* method may use information in its property structure, the state of its ports, and its requested local and global/persistent member data to decide what to do. Normally this involves using messages in buffers at input ports to produce messages in buffers at output ports.

The `timedOut` input argument indicates whether the *run* method is being invoked due to time passing (the `usecs` value of the run condition). This indication is independent of port readiness other than indicating that the run condition was *not* met.

C Language: Each port's `current.data` member may be tested to indicate port readiness if that port is *not* included in the run condition.

C++ Language: Each port is a member data object whose name is the name of the port. The `hasBuffer()` method on that port object (e.g. `indata.hasBuffer()`), returns whether the port is ready.

Normally this test is unnecessary since the container only calls this *run* method when the run condition is satisfied, and that implies that all ports included in the run condition have buffers (i.e. are *ready*).

The *run* method can indicate that all ports should be advanced by the special return value `RCC_ADVANCE`. It can also indicate disposition of buffers and ports by using the **release**, **send**, **request**, **take**, or **advance** container methods.

C Language: The *run* method may change the run condition by setting a TRUE value to the location indicated by the `newRunCondition` output argument, and setting a pointer to a new run condition in the `runCondition` member of `RCCWorker`. The `runCondition` member of `RCCWorker` is initially set to the value in the `RCCDispatch` structure (including `NULL` to indicate the default run condition).

C++ Language: The *run* method may change the run condition using the `setRunCondition` container method passing a pointer to a (not *const*) new run condition object. The current run condition (pointer) is retrieved using the `getRunCondition()` method. The run condition for a worker is initially set to the default, which is retrieved as `NULL`.

If the `runCondition` is set to `NULL`, the default run condition is restored. The run condition for a worker is initially set to the default. Since the worker code is managing the run conditions (passed by pointer), it can use it as a convenient “state variable” when its execution modes each have a different run condition.

Code in the *run* method accesses information about ports and current buffers by accessing objects/structures that are ports in the Worker object.

C Language: The *run* method accesses information about current buffers through members of each port's `RCCPort` structure member in the `ports` array in the `RCCWorker` structure (e.g. `self->ports[n]`, where `n` is the ordinal for the port). The `current.data` member is the pointer to the message data for both input and output ports. For input ports, the `input.length` structure member of the `RCCPort` is the length of bytes of the message in the current buffer, and `input.u.operation` is the opcode for the current input message. For output ports, the `output.length` is the length of bytes of the message in the current buffer, and `output.u.operation` is the opcode for the current output message.

C++ Language: The *run* method accesses port status and buffer contents using the worker's object port data member methods.

If a port is in the worker's run condition, then it can assume a current buffer is present at that port when the run method is entered. If not, the worker can test whether there is a current buffer; in C, the port's `current.data` member is non-`NULL`, in C++, the port's `hasBuffer()` returns true.

The full definition of buffer and port object/structure types is in the container method section below.

4.1.9.3 Returns

This method shall return a `RCCResult` value.

The value **RCC_ADVANCE** indicates that all ports should be advanced that were ready on entry to the *run* function and were not subject to container methods called since then (e.g. *advance*, *request* etc. described below).

If the *run* method cannot succeed, it should return **RCC_ERROR**, indicating that it should not be called again (regardless of run condition). If the worker detects an error that would disable the implementation or its environment, it should return **RCC_FATAL**. Otherwise it should return **RCC_OK** or **RCC_ADVANCE** if normal worker execution should proceed (the *run* method will be called again when the run condition is true).

The run method should return **RCC_DONE** if no further worker execution should happen. Returning **RCC_DONE** indicates to the container that the worker will never require further execution. Providing this advice to the container allows it to take advantage of this fact and possibly release resources (such as I/O buffers) prior to the worker instance being destroyed or reused.

Returning **RCC_ADVANCE_DONE** combines the meaning of **RCC_ADVANCE** and **RCC_DONE**.

4.2 Container methods, called by the worker.

These methods are those that the worker calls to invoke some functionality provided by the container. The container methods are in three categories: container scope (this section), port scope (next section) and buffer scope (following section).

The use of all container methods is optional, and typically unneeded for simple workers that only use run conditions and the *run* method. They are mostly used to provide additional flexibility and functionality in message handling.

The methods are described below, and the following section details the data types used by the arguments of the container methods. **All container methods are non-blocking.**

C Language: container methods are accessed using function pointer members in the container structure member of the RCCWorker structure, e.g.:

```
self->container.setError(...)
```

C++ Language: container methods exist in an accessible base class for the worker and are thus called directly. Container methods relating to ports and buffers are in fact methods of the worker's member data objects for ports, rather than on the worker (or container) objects.

4.2.1 *time/getTime* — container method

This method is used by the worker to retrieve the time of day as a 64 bit unsigned number that represents GPS time in units of 2^{-32} seconds (~233 ps). Thus the most significant 32 bits represent GPS time in seconds. The accuracy of this time is dependent on the container implementation.

Note that GPS time is monotonic, as opposed to UTC time (the POSIX standard), which is subject to leap seconds etc.

4.2.1.1 Synopsis

```
RCCTime (*time)(); // C Language  
RCCTime getTime(); // C++ Language
```

4.2.1.2 Returns

This function returns GPS time in units of 2^{-32} seconds.

4.2.2 *setError* — container method

This method is used by the worker to convey an error message associated with returning `RCC_ERROR` or `RCC_FATAL`. It has semantics similar to `printf`.

4.2.2.1 Synopsis

```
RCCResult (*setError)(const char *fmt, ...); // C Language  
RCCResult setError(const char *fmt, ...); // C++ Language
```

4.2.2.2 Returns

This function returns `RCC_ERROR`, which allows a worker to simultaneously set the error and return the error indication. For example, a worker method could have (in C):

```
if (some error condition)
    return self->container.setError("Failure due to %d", arg);
```

Or, in C++:

```
if (some error condition)
    return setError("Failure due to %d", arg);
```

If the worker must be return `RCC_FATAL`, it must call this function separate from returning that value.

4.2.3 **getRunCondition** — container method — C++ Language only

This method is used by the worker to retrieve its current run condition. If the current run condition is specified to be the default, `NULL` is returned. Otherwise the pointer returned is the most recent one passed from the worker to the container in the `setRunCondition` method.

4.2.3.1 Synopsis

```
RCC::RunCondition *getRunCondition();
```

4.2.3.2 Returns

This function returns a pointer to the run condition most recently set by the worker using the `setRunCondition` call, or `NULL`.

4.2.4 **setRunCondition** — container method — C++ Language only

This container method is used by the worker to set a new run condition. If the run condition supplied is `NULL`, the default run condition becomes active, which is to wait for all ports to be ready, with no timeout. If the worker needs to start execution with a non-default run condition, it can call this container method in its constructor, or in its `initialize` or `start` methods.

4.2.4.1 Synopsis

```
void setRunCondition(RCC::RunCondition *rc);
```

4.2.4.2 Returns

No return value.

4.2.5 The **RCC::RunCondition** C++ class

The `RCC::RunCondition` type represents the run conditions described above, and holds the information used by the container to determine when it is appropriate to invoke the `run` operation of a worker. It is a C++ Struct, with convenience constructors. The defined members are always written by the worker, and never by the container. The following table describes the accessible members of the `RunCondition` structure.

Table 6: RunCondition Members

Member Name	Member Data Type	Member Description
portMasks	RCCPortMask *	A pointer to a zero-terminated array of port masks, each of which indicates a bit-mask of port readiness. The run condition is considered <i>true</i> when any of the masks is true. A mask is <i>true</i> when all indicated ports are ready (logical AND of port readiness). A port is indicated by its bit being set ($1 \ll \text{port_ordinal}$). If the pointer itself is NULL , the run condition is always true. A bit set for an unconnected port is ignored, thus the default run condition can be used with unconnected ports.
timeout	bool	Indicates that the usecs member determines when enough time has passed to make the run condition true. This value is used to enable or disable the timeout, without changing usecs .
usecs	uint32_t	If this amount of time has passed (in microseconds) since the run method was last <i>entered</i> , the run condition is true.

Thus the overall run condition is the logical OR of the **portMasks** and the timeout. If the worker offers no **run** method in its RCCDispatch structure (see below), run conditions are ignored. If the **portMasks** member is **NULL**, it indicates that no port readiness check is performed and thus the run condition is always true. Typical combinations are:

Table 6: Run Condition Combinations

Shorthand	Port mask	Timeout	RunCondition Description
Always run	portMasks == NULL	ignored	the worker is always ready, no timeout
Run when data ports are ready	portMasks != NULL	False	Run condition is TRUE when any mask is true. If there are no masks (portMasks[0] == 0), then the run condition is always false.
Run when data ports are ready or timeout.	portMasks != NULL and portMasks[0] != 0	True	Run condition is TRUE when any mask is true OR if the timeout expires. The timeout will take effect if time passes without the masks being satisfied.
Periodic execution	portMasks != NULL and portMasks[0] == 0	True	Since portMasks[0] == 0, port masks can never be true, thus this establishes period execution independent of port readiness.

The RunCondition structure is defined with public members as defined above, and several convenience constructors:

Table 6: RunCondition Constructors

<code>RunCondition()</code>	Specify the default run condition, no timeout, all connected ports must be ready
<code>RunCondition (RCC::PortMask first, ...)</code>	Specify a list of port masks as variable arguments that are a zero terminated list of masks
<code>RunCondition (RCC::PortMask *masks, uint32_t usecs = 0, bool timeout = false)</code>	Specify all fields, with defaults for timeout. RunCondition(NULL) indicates: <i>always run</i> .

4.3 Port Management Data Members and Methods

This section describes port management methods, whereas the next section describes some members and methods used to separately manipulate buffers. The current buffer at a port is in fact part of the port object, so methods on the current buffer are in fact methods on the port. In the cases where buffer methods are used when the buffer is *not* the current buffer at some port, specific buffer methods are used, and those are described in the next section.

In the C language, port management methods are directly accessed function pointer members in the `RCCWorker` data structure, like all container methods. Each has a `RCCPort` pointer as an argument.

In C++, each port is a member data object (of the worker object) whose name is the name of port (from the OCS). Thus port management methods are methods on that object.

As an example, to *advance* a port in C, and optionally access the data of the now-current buffer, you would write:

```
RCCPort *inport = &self->ports[MY_IN_PORT];

if (self->container.advance(inport)) {
    void *p = inport->current.data;
    ...
}
```

In C++ this would be:

```
if (inport.advance()) {
    void *p = inport.data();
    ...
}
```

Note that in C++, the current buffer object is *inherited* by the port object, so the port data member is the object used to access the current buffer's `data()` method.

4.3.1 *advance* — port method (C++), container method (C)

This method releases the current buffer at a port, *and* requests that a new buffer be made available as the current buffer on the port. This is a convenience/efficiency combination of the **release** (the current buffer) and **request** (a new current buffer) methods described below.

An optional minimum size in bytes may be requested (0 is the default in C++, 0 can be supplied in C for no minimum size).

4.3.1.1 Synopsis

```
RCCBoolean (*advance)(RCCPort *port, size_t minSize); // C language
bool Port::advance(size_t minSize = 0);                // C++
```

4.3.1.2 Returns

A boolean value is returned indicating whether the **request** was immediately satisfied.

4.3.2 **hasBuffer** — *query a port for whether it has a current buffer-to-release*

This method returns whether a port has a current buffer. If a port is already part of a Worker's run condition, and it is in every port mask of the run condition, it can be assumed to have a current buffer whenever the worker is run. This is the case for 90% of all workers. Thus this method is only needed and used when a port is not in the run condition, as might be the case for a port that received some exceptional condition message.

In C, this is a port structure member access. In C++ it is a port method.

4.3.2.1 Synopsis

```
Port->current.data != NULL          // C Language
bool RCCUserPort::hasBuffer();      // C++ Language
```

4.3.2.2 Returns

For C++, a bool value is returned indicating whether the port has a current buffer.

4.3.3 **isConnected** — *query a port for being connected*

This method returns whether a port is currently connected or not. Component specifications (OCS) indicate whether each port of a component (and thus all workers) must be connected before workers are started or run. A port is considered “optional” for a worker if the worker can tolerate the port *not* being connected. If a port is *not* optional, the worker code can assume it is connected, and this method would not be needed or used.

In C, this is a Worker (self) structure member access. In C++ it is a port method.

4.3.3.1 Synopsis

```
self->connectedPorts & (1 << PortOrdinal); // C Language
bool RCCUserPort::isConnected();           // C++ Language
```

4.3.3.2 Returns

For C++, a bool value is returned indicating whether the port is connected.

4.3.4 **request** — *port method (C++), container method (C)*

This method *requests* that a new buffer be made available as the current buffer on a port. If the port already has a current buffer, the request is considered satisfied. This request indicates to the container that it should make a buffer available when possible. Without an explicit request, the container may not make a buffer available since that would dedicate resources when they may not be needed. An implicit **request** is made for all ports that are part of the current run condition.

An optional minimum size may be requested.

Note that this method is not used or needed when ports are *advanced*, only when buffers are being explicitly managed (e.g. released, etc.).

4.3.4.1 Synopsis

```
RCCBoolean (*request)(RCCPort *port, size_t minSize); // C Language
bool Port(size_t minSize = 0);                       //C++ language
```

4.3.4.2 Returns

A boolean value is returned indicating whether the request was immediately satisfied.

4.3.5 **send** — *send an input buffer on an output port*

This method sends an input buffer on an *output* port. If the buffer is a current buffer for a port, this buffer will no longer be that port's current buffer. Buffer ownership passes back to the container. This method is used to effect zero-copy transfer of a message from an input port to an output port.

C Language: The operation and message length are supplied as arguments.

C++ Language: The operation and message length are attributes of the buffer object and can be changed before it is sent using the `setOpCode` and/or `setLength` methods.

The sent buffer is either a current buffer of an input port or a buffer taken from an input port.

4.3.5.1 Synopsis

```
void (*send)(RCCPort* port,
             RCCBuffer* buffer,
             RCCOrdinal op,
             uint32_t length); // C Language
void RCCPort::send(RCC::Buffer &buffer); // C++ Language
```

4.3.5.2 Returns

Nothing.

4.3.6 **setDefaultLength** — *set the default message length at an output port*

This method sets the default message length in bytes for a port. It is useful when messages produced at a port will always be the same size.

This method is used only when the message content is manipulated as a raw untyped buffer, rather than using the message content access methods described below.

The initial setting of the default message length is the length of the buffer.

This method is only available in C++.

4.3.6.1 Synopsis

```
void RCCUserPort::setDefaultLength(size_t length); // C++ Language
```

4.3.6.2 Returns

Nothing.

4.3.7 **setDefaultOpCode** — *set the default opcode at an output port*

This method sets the default opcode for an output port. It is useful when messages produced at a port will always have the same opcode.

The initial setting of the default opcode is zero (indicating the first operation specified in the associated protocol).

This method is only available in C++.

4.3.7.1 Synopsis

```
void RCCUserPort::setDefaultOpCode(RCCOpCode opcode); // C++ only
```

4.3.7.2 Returns

Nothing.

4.3.8 **take** — *port method (C++), container method (C)*

This method *takes* the current buffer from a port, and optionally releases a previously taken buffer. This method is used when workers need to maintain a history of one or more previous buffers while still requesting new buffers, e.g. for “sliding window” algorithms.

The optional buffer to be released is provided by a possibly **NULL** pointer to a buffer. This buffer-to-release argument is a convenience feature of the API to allow cycling of buffers in a single call, e.g.:

```
class XyzWorker {
    RCCUserBuffer *m_prev;           // remember buff from previous run
    XyzWorker() : m_prev(NULL) {}
    void process_curr_and_prev(RCCUserBuffer *curr,
                              RCCUserBuffer *prev) {}

    run() {
        process_curr_and_prev(&in, m_prev);
        m_prev = &in.take(m_prev);
    }
}
```

In C, a container method has an **RCCPort** pointer and a possibly **NULL** buffer-to-release pointer as inputs, and a buffer pointer (thus a structure provided by the caller) where the taken buffer structure will be copied (not the data, just the structure).

In C++, **take** is a port method and the taken buffer is simply returned via a reference.

Ownership of the taken buffer is passed to the worker. The current buffer now taken is no longer the current buffer. This method is used when the worker needs access to more than one buffer at a time from an input port. **Take** implies **request** (to get another current buffer).

It is an error to call this method when the port does not have a current buffer.

4.3.8.1 Synopsis

```
void (*take)
(RCCPort* port, RCCBuffer* releaseBuffer,
 RCCBuffer* takenBuffer);           // C language
RCCUserBuffer &Port::take
(RCCBuffer *release = NULL);        // C++ Language
```

4.3.8.2 Returns

Nothing is returned in C, since the taken buffer object is copied to the location supplied by the **takenBuffer** argument. In C++, a reference to the taken buffer is returned.

4.4 Buffer Management Data Members and Methods

This section describes methods that apply to buffers separate from ports. As mentioned above, when a buffer is the current buffer at a port, the port object itself is normally used with these buffer methods. In C++, the port object inherits the buffer object that is the current buffer. In C, the current buffer is the `current` member of the `RCCPort` structure.

When methods deal with “opCodes”, they are dealing with an ordinal for the message in a buffer, which identifies which type of message is in the buffer, among those defined by the protocol. OpCode values are zero-origin in the order of how the operations are defined in the protocol (OPS) XML file. Enumeration constants for the messages in a protocol are generated and are described in the code generation section below.

4.4.1 *checkLength* — check size of the buffer

Check that a message of a given size in bytes will fit into the buffer, usually associated with an output port. An exception is thrown if the message will not fit. This is useful when the worker is creating a variable length message and wants to ensure it will fit into the available buffer. This check results in an error that the worker cannot check, and is intended to avoid unexpected buffer size mismatches, similar to the “assert” standard library function.

This method is used only when the message content is manipulated as a raw untyped buffer, rather than using the message content access methods described below.

This method is only available in C++.

4.4.1.1 Synopsis

```
void RCCUserPort::checkLength(size_t neededSize); // C++ Language
```

4.4.1.2 Returns

Nothing.

4.4.2 *data* — access the raw contents of the buffer

This method returns a pointer to the raw data in a buffer. It is analogous to the “data” method in C++ STL container classes, although it has no type.

4.4.2.1 Synopsis

```
void *p = buffer->data;           // C Language  
void *RCCUserBuffer::data();      // C++ Language
```

4.4.2.2 Returns

The value returned is a void pointer to the contents of the buffer.

4.4.3 **length** — retrieve the length of the message in a buffer – C++ only

This method returns the number of bytes of the message in the buffer, which is typically an input buffer. This is C++ only. When using C, access to the length is via the current buffer at a port, e.g. `port->input.length`.

4.4.3.1 Synopsis

```
size_t RCCUserBuffer::length() const;    // C++ Language
```

4.4.3.2 Returns

The method returns the length in bytes of the message in the buffer.

4.4.4 **maxLength** — retrieve the maximum available space in the buffer.

This method allows the worker to retrieve the actual size of the buffer, typically used on output buffers.

4.4.4.1 Synopsis

```
size_t len = buffer->maxLength;          // C language  
size_t RCCUserBuffer::maxLength() const;  // C++ language
```

4.4.4.2 Returns

The size of the buffer in bytes is returned.

4.4.5 **opCode** — retrieve the opCode of the message in a buffer — C++ only

This method retrieves the *opCode* of the message in the buffer, typically an input buffer. This is C++ only. When using C, access to the opCode is via the current buffer at a port, e.g. `port->input.u.operation`

4.4.5.1 Synopsis

```
RCCOpCode RCCUserBuffer::opCode() const;    // C++ Language
```

4.4.5.2 Returns

This method returns the opCode of the message in the buffer.

4.4.6 **release** — release a buffer

This method releases a buffer for reuse. If the buffer is the current buffer for a port, it will no longer be the current buffer. Buffer ownership passes back to the container. Buffers for a port **must be released in the order obtained**, per port. Note that this method is not used or needed when ports are “advanced”, only when buffers are obtained from a port using other port management functions such as **take**.

Releasing a current buffer does not imply requesting a new current buffer: that request must be explicit. A release without a request might be useful if a worker is entering a mode where it no longer needs any more data from an input port or it no longer needs to send any more buffers on an output port.

In C, this is a container method. In C++ it is a buffer method.

4.4.6.1 Synopsis

```
void (*release)(RCCBuffer* buffer); // C Language
void RCCUserBuffer::release();      // C++ Language
```

4.4.6.2 Returns

Nothing.

4.4.7 **setInfo** — set the metadata associated with a buffer — C++ only

This method is a convenient combination of setting both the **opCode** and the **length** of the message in a buffer. It is typically used on output buffers, when both opCode and length are being set at the same time.

4.4.7.1 Synopsis

```
void RCCUserBuffer::setInfo(RCCOpCode op, size_t length); // C++
```

4.4.7.2 Returns

Nothing.

4.4.8 **setLength** — set the length of the message in a buffer — C++ only

This method sets the length in bytes of the message in a buffer.

This is C++ only. When using C, access to the length of an output buffer is via the current buffer at a port, e.g. `port->output.u.length`, or via the **send** port method.

4.4.8.1 Synopsis

```
void RCCUserBuffer::setLength(size_t length); // C++ Language
```

4.4.8.2 Returns

Nothing.

4.4.9 **setOpCode** — set the opCode for the message in a buffer — C++ only

This method sets the opCode of the message in a buffer.

This is C++ only. When using C, access to the opCode of an output buffer is via the current buffer at a port, e.g. `port->output.u.operation`, or via the **send** port method.

4.4.9.1 Synopsis

```
void RCCUserBuffer::setOpCode(RCCOpCode op); // C++ Language
```

4.4.9.2 Returns

Nothing.

4.4.10 *topLength* — *retrieve the size of the single sequence in a message* — C++ only

This method only applies to input buffers that contain messages consisting of a single sequence argument (as defined in the protocol). It retrieves the length, *in elements*, of that sequence. The **elemSize** argument to **topLength** is the size in bytes of the elements of the sequence in the buffer. An error check is made to ensure that the size of the message in bytes is divisible by the **elemSize** value provided.

Normally, sequence lengths are retrieved using the message/operation access methods described below in the “accessing the contents of messages” section. However, when message contents are accessed as an untyped raw buffer, this method allows the worker to retrieve the number of elements of a given size in the message without the container knowing the data type of the message. It performs the size conversion and error check in one API.

The message must consist of exactly one sequence for this method to be valid. Sequences can have zero elements.

This method is only available in C++.

4.4.10.1 Synopsis

```
size_t RCCUserBuffer::topLength(size_t elemSize); // C++ Language
```

4.4.10.2 Returns

For C++, this method returns a `size_t` value indicating the number of sequence elements in the message, given the size of elements known by the worker.

4.5 Accessing the Contents of Messages

The sections above described methods used to access the content and metadata (length and opCode) of messages in buffers. Those methods operate independent of the actual types of the data in the messages, and put the burden of accessing typed data on the worker code. I.e. the worker code would have to convert (cast) the pointer types and manually deal with the multiple data types for the arguments in a message.

Accessing the individual arguments (fields) in the payload for an input message (or setting the fields of an output message) is facilitated differently in C vs. C++.

In C, a data structure is generated that can be overlaid on the input buffer to access all fields up to and including the first variable length field (sequence or string). In fact, for each port, a union of data structures is defined for all the possible messages that may be in a buffer for a port. These structures are described in the code generation section below, but the basic pattern is, for a port PXY, with protocol PR1 with operations Op1 and Op2, there is a union of structures like this:

```
union PxyOperations {
    struct Op1 {
        ... fields in the message for Op1 messages ...
    } op1;
    struct Op2 {
        ... fields in the message for Op2 messages ...
    } op2;
};
```

Note that the data types are camel-cased, while the union members are lower cased. So, in C you can access messages based on opCode like this:

```
union PxyOperations *up = port->current.data;
switch (port->input.u.operation) {
case Pr1_Op1:
    use up->op1 as message structure; break;
case Pr1_Op2:
    use up->op2 as message structure; break;
}
```

The major limitation of this mechanism for C language workers is that the structures only cover the fields of messages up to and including the first variable length field. After that, the worker code must manually access the remaining fields using pointer arithmetic and casting. This covers most protocols used in simple systems, but is painful when the limitation is exceeded.

For C++ workers, it is easier and does not have the same limitations. Each port is a data member of the worker object whose name is the port's name in the OCS, and each such port data member has an accessor for the information for each possible message in the port's protocol. Finally, each field of the payload for each operation also has an accessor. Thus, for the example above, accessing the Arg1 field of the Op1 operation on the protocol for port "in", would simply be:

```
in.Op1().Arg1()
```

These C++ message access methods are `const` and return a value for input ports, and return a non-`const` reference for output ports.

When the argument type is a sequence or array, the field accessor returns a reference to an object that has two methods that act just like the C++ STL `vector` collection type:

```
<arg-data-type> *data() const;
size_t size() const;
```

Thus if “Arg1” in the above example was a sequence of unsigned shorts, you could access this field with:

```
uint16_t *vals = in.Op1().Arg1().data();
size_t nvals = in.Op1().Arg1().size();
```

For output ports, the `data()` accessor for sequences or arrays is not `const`, and there are two other methods that also act like the C++ STL vector type:

```
<arg-data-type> *data();
size_t capacity() const;
void resize(size_t size);
```

4.6 Property Access Notifications

Workers sometimes need dynamic notifications when properties are written or read by control software. i.e. the worker would like to run some code just before a property read returns (to potentially create the correct value to read) or after a property write (to cause some other side effect of the new written value).

These capabilities are enabled when the property in the OWD (not OCS) has the `readSync` or `writeSync` attributes set to true. The `readSync` attribute being true indicates that the worker would like to run some code *before* the value of the property is passed back to the control software reading the property. The `writeSync` attribute being true indicates that the worker would like to run some code *after* a new value is written, in order to implement some side effect when the property value is changed.

An example of using `readSync` is when the value should be computed based on some other dynamic condition, e.g. reading a real-time or physical sensor value. An example of `writeSync` is when writing a new value should atomically affect some other state of the worker.

Using the C language, the notifications are made via the calls to the `beforeQuery` (for any `readSync` property) and `afterConfigure` (for any `writeSync`) control methods. The worker cannot know which properties were written or read, but at least knows that only `readSync` or `writeSync` properties could have caused these notifications.

In C++, when these attributes are set, the skeleton includes empty implementations of notification methods for each such property that the worker author can fill out. For `readSync` properties a worker class member function called `<property>_written` must be present. For `writeSync` properties, a member function named `<property>_read` must be present. Both notification member functions take no arguments and return `RCCResult`, which allows these functions to indicate errors. A `readSync` notification function can set the property value locally and then return, and

that value will then be conveyed back to control software. A **writeSync** notification is called after the new value has been written, so the function can locally access the new value when it is called.

4.7 Controlling Slave Workers from Proxies (C++ only)

When the RCC OWD indicates that an RCC worker is a proxy for another (slave) worker, it gives the proxy worker convenient access to the slave's control interface (for control operations and property accesses). The primary reason for proxy workers is to standardize and/or make more user-friendly, the control of the underlying slave worker. The slave worker is frequently device-specific, having esoteric properties or configuration requirements that normal users should not be concerned with.

As a concrete example, a “device worker” that controls an RF front-end device typically wants to expose the native functionality of the device in a high performance and/or resource-conservative fashion. However, most application users would only want a subset of those options that are uniform across all RF front-ends they might use.

Delegating the normalization of the control/configuration interface to the RCC proxy worker can expose to users a high level configuration interface with complex property data types as required. This relieves this burden from the embedded device worker which may be in an environment (e.g. on a small FPGA) that makes this normalization difficult and/or expensive. Furthermore, this enables the device worker to faithfully and simply implement the device's native control interface for direct use by lower level applications that want to exploit lower level device-specific features, without cluttering up the proxy's more generic interface.

A proxy worker can provide a higher level control interface, and also can act as the software module that encapsulates specialized configuration and control sequencing of the device, without requiring the application to use special APIs in a custom application. The proxy worker is “just another component” in the application that is fully described in XML.

Inside the code of a proxy, the slave worker is accessed by using the **slave** data member of the worker's class object. This data member has member functions for the control operations on the slave (except **initialize** and **release**). For each readable property of the slave worker, a **get_<property-name>** member function is available to retrieve the current value of the slave's property, and for each writable property of the slave worker a **set_<property-name>** is available. If the property is an array, both the **get** and **set** member functions for that property (of the **slave** data member) have arguments specifying the index in each dimension. The **set** functions then have a value argument for the new value of the property. Here are some examples, usable in any proxy worker member function (e.g. **start**, or **run**):

```
uint8_t x = slave.get_byteregA();
slave.set_byteregA(3);
x = slave.get_arrayprop(2);
slave.set_arrayprop(3, x);
x = slave.get_array2d(1, 2);
slave.set_array2d(1, 2, x);
slave.stop();
slave.start();
```

There is currently no way to get or set whole array values, only individual elements.

4.8 Worker Dispatch Structures — C Language only

When workers are loaded for execution by the container, the container finds the worker by getting access to its `RCCDispatch` structure. Thus the only external symbol the worker code needs to define is the symbol that holds this structure.

The code generator generates this initialization in the skeleton file, and it can be further customized by the worker author.

4.8.1 *RCCDispatch Structure Type*

This type is the “dispatch table” for the operations of the C-language worker interface. It represents the functionality a worker provides to a container when it is loaded. The container must gain access to this structure when the worker is loaded and executed. All members are statically initialized by the worker source code. This structure also contains other descriptive information required by the container to use the worker.

The RCC C-language skeleton will contain a default initialization for this structure based on what is in the OWD, which looks like (for worker `xyz.rcc`):

```
XYZ_METHOD_DECLARATIONS;
RCCDispatch bias = {
    /* insert any custom initializations here */
    XYZ_DISPATCH
};
```

Any non-default member initializations must be placed after the `XYZ_DISPATCH` line, and be specified using the named member syntax, e.g.:

```
RCCDispatch bias = {
    /* insert any custom initializations here */
    XYZ_DISPATCH
    .memSize = 16*sizeof(MyData),
    .runCondition = &myRunCondition
};
```

The members of this structure that may be set this way are defined in the following table:

Table 6: Members of the RCCDispatch structure — C-language only

RCCDispatch Member	Member Data Type	Member Description
memSizes	uint32_t*	This zero-terminated array of memory sizes indicates allocations required by the worker. Multiple allocations allow the worker to avoid aggregating its requirements in a single allocation. May be NULL, when no allocations are required.
runCondition	RCCRunCondition*	The initial run condition used. If this pointer is NULL it implies a run condition of all connected ports being ready and no timeout. If there are no ports, the default is no port masks to check, indicating always ready to run.
optionalPorts	RCCPortMask	A mask indicating ports that may be unconnected. The default, 0, means all ports must be connected before the worker is started.
memSize	size_t	This size indicates a required memory allocation. May be zero, indicating no allocation is required.

5 Code Generation for RCC Workers

The worker's OCS and OWD are used to generate two code files, the header and the skeleton. The header file, which should not be edited, contains type definitions customized for the worker. This file is named `<workername>_Worker.h` (in the C language) or `<workername>-worker.hh` (in the C++ language), *without* capitalizing the worker name. It is placed in the “**gen**” subdirectory of the worker's directory.

The skeleton file is a small file that is generated as the basis for writing the worker's execution code (e.g. the run method implementation). It relies heavily on the generated header file and thus is as small as possible to avoid code changes when the OCS or OWD is changed. The skeleton is generated in the **gen** subdirectory with the name `<workername>-skel.c` (C language) or `<workername>-skel.cc` (C++ language). It is also copied into the worker's directory as `<workername>.c` or `<workername>.cc` and *this* copy is what should be edited to add the worker's functional code. The `<workername>-skel.c` file should not be edited, but left for reference to see exactly what the code generator produced.

In the descriptions below, names in angle brackets are names specific to the worker, as found in the worker's OCS and OWD XML files.

5.1 *Namespace Management*

To avoid name space collisions, some rules are used by the code generator when creating the header and skeleton files.

In the C++ language, generated data types in the header are placed in a namespace whose name is: `<worker>WorkerTypes`. In the actual worker source file, this namespace is normally imported, via the “`using namespace`” directive. The actual worker derived class is a class in the global namespace whose name is `<worker>Worker`. The dispatch entry point, which is necessarily a C external symbol, has the worker's name as its external symbol.

In the C language, the worker's dispatch table is given the external symbol name of the worker, and all methods are declared `static`. All data types use the capitalized worker name as a prefix and all constants and macros use the upper cased worker name followed by an underscore as a prefix.

5.2 Generated Data Types

Various data types are generated in the header file for the worker. They are :

- Port name enumeration
- Property data structure
- Opcode enumeration for each protocol used
- Opcode enumeration for each port (C only)
- Message structures (C only)
- Worker base class (C++ only)

5.2.1 The enumeration constants for the worker's ports

An enumeration type is generated in the header file which defines constants for the ordinal of each port, of the form `<worker>_<port>`, all upper case. These ordinals can be used when creating port masks for run conditions, or, for C workers, indexing into the port array in the `RCCWorker` structure.

5.2.2 The properties structure type

This is the generated structure definition that reflects the properties declared for the worker in its OCS and OWD XML files. Each member of this structure has the data type corresponding to the property's description in the OWD and OCS XML files. Members for properties that are *not* declared `volatile` in XML are `const` in this structure, indicating that the worker is not expected to change their values.

The standard name of the struct type (typedef name in C) is `XyzProperties` where "Xyz" is the *capitalized* worker implementation name.

Properties which are sequences are generated as a structure containing a `uint32_t` member whose name is `length`, holding the number of elements in the sequence. Padding may be added before and after the length member to achieve the required alignment of this length field as well as the sequence data following it. The member name of the structure is the property name. The members of the structure are `length` and `data`. The data member is a C array whose length is the `SequenceLength` attribute from the Property element in the XML. Remember that for properties (as opposed to protocols), sequences and strings must have a maximum size (i.e. be bounded).

Struct properties have structure tags the same name as the property name, preceded by the worker name.

The names of the integer types for properties: `short`, `ushort`, `long`, `ulong`, `longlong`, `ulonglong`, `uchar`, `char` are mapped to the respective `<stdint.h>` types: `int16_t`, `uint16_t`, `int32_t`, `uint32_t`, `int64_t`, `uint64_t`, `uint8_t`. The OCS/OWD `char` type is mapped to the `RCCchar` type, which is normally mapped to the basic C/C++ `char` type and it is assumed to be signed.

The names of the non-integer types for simple properties are the OCS-defined type names capitalized and prefixed with “RCC”: e.g. **RCCBoolean**, **RCCChar**, **RCCFloat**, **RCCDouble**.

Properties that are string properties are **RCCChar** arrays whose size is one more than the **StringLength** attribute of that string property in the Property element, and the values are null terminated strings.

5.2.3 Structures for Message Payloads (C only)

For each port a union type is defined for all possible messages at that port. Members of the union type are generated structures for each possible message in that protocol. The name of the union type is **<Port>Operations**, with port name capitalized.

For each message type specified in the protocol, a structure is defined whose type name is the capitalized name of the operation in the protocol, and whose member name is the lower cased name. Arguments to the operation in the protocol are structure members of that per-operation structure. The structure layout is as defined above for the property structure, with the exception that variable size elements (strings and sequences) are allowed, and are sized as [1], with no further members generated after the first variable sized member.

The structure generated is padded and packed. This means padding members are explicitly inserted to ensure all types are aligned on their own boundaries. It also means that the compiler is told to “pack” the structure. Given the insertion of padding, the “packing” simply means that there is no padding at the end of the structure to achieve any alignment. Thus the struct definition should not be used in any other context where this lack of overall size alignment is required (i.e. an array of such structures).

There is one special exception for messages that consist entirely of one sequence of fixed length elements: the length of the sequence is implied by the overall message length as specified by the “length” value of buffers at input and output ports, and *not* represented by inserting a **uint32_t** value in the message buffer/structure. In all other cases (not-fixed sequence elements, or more than one top level data argument in the message), sequences in messages are represented by a struct containing **length** and **data** members.

For example, for a port named “in”, with a protocol whose first operation was “sampledata”, and whose first argument was “sample” - a sequence of **Ulonglong** values, the worker’s custom port union would look like:

```
union InOperations {
    // Structure for the 'sampledata' operation on port 'in'
    struct __attribute__((__packed__)) Sampledata {
        uint64_t sample[1];
    } sampledata;
};
```

5.2.4 *Worker Base Class (C++ only)*

For C++ workers, a base class is generated in the header file that is inherited by a derived class in the skeleton. The base class is generated to maximize the code generated in the header file and minimize the code generated in the skeleton.

The worker base class is generated to support all the documented features described above, including access to properties, and contain pure virtual member function declarations for all the member functions required to be implemented in the derived class. The details of the code generation are implementation defined and subject to change.

6 RCC Local Services

Local service APIs support RCC workers running in a container. Containers are required to supply them, and workers are constrained to use only them (other than the worker-to-container entry points defined above). The RCC local services are defined as a small subset of the POSIX and ISO-C runtime libraries.

These local services APIs represent a minimal environment required of embedded systems. If other APIs are used, there is no guarantee they will be available in all containers.

6.1 RCC Local Services AEP as a small subset of POSIX and ISO-C.

The local services available for RCC workers are typical runtime support functions without the functions that require heavy weight OS support. All I/O is excluded since portable RCC workers should be performing all I/O via the OpenCPI data plane ports. Some containers in fact allow many more functions to be called from a worker, but that makes such workers non-portable.

The POSIX Minimal Realtime System Profile (PSE51) from the IEEE Std 1003.13TM-2003 standard is used as the superset of functionality that is reduced (with subtractions) to the RCC local services available functions. Using this subset requires setting the `_POSIX_AEP_RT_MINIMAL_C_SOURCE` feature test macro to the value 200312L during compilation.

The PSE51 specification defines POSIX.1 units of functionality in its table 6-1. This RCC Local Services AEP (application environment profile) removes these units of functionality from that table:

- `POSIX_DEVICE_IO`
- `POSIX_FILE_LOCKING`
- `POSIX_SIGNALS`
- `XSI_THREAD_MUTEX_EXT`
- `XSI_THREADS_EXT`
- `POSIX_THREADS_BASE`

The retained units of functionality, and their defined symbols and functions are:

POSIX_C_LANG_JUMP: `longjmp()`, `setjmp()`;

POSIX_C_LANG_SUPPORT: `abs()`, `asctime()`, `asctime_r()`, `atof()`, `atoi()`, `atol()`, `atoll()`, `bsearch()`, `calloc()`, `ctime()`, `ctime_r()`, `difftime()`, `div()`, `feclearexcept()`, `fegetenv()`, `fegetexceptflag()`, `fegetround()`, `fehldexcept()`, `feraiseexcept()`, `fesetenv()`, `fesetexceptflag()`, `fesetround()`, `fetestexcept()`, `feupdateenv()`, `free()`, `gmtime()`, `gmtime_r()`, `imaxabs()`, `imaxdiv()`, `isalnum()`, `isalpha()`, `isblank()`, `iscntrl()`, `isdigit()`, `isgraph()`, `islower()`, `isprint()`, `ispunct()`, `isspace()`, `isupper()`, `isxdigit()`, `labs()`, `ldiv()`, `llabs()`, `lldiv()`, `localeconv()`, `localtime()`, `localtime_r()`, `malloc()`, `memchr()`, `memcmp()`, `memcpy()`, `memmove()`, `memset()`, `mktime()`, `qsort()`, `rand()`, `rand_r()`, `realloc()`, `setlocale()`, `snprintf()`, `sprintf()`, `srand()`, `sscanf()`, `strcat()`, `strchr()`, `strcmp()`, `strcoll()`, `strcpy()`, `strcspn()`, `strerror()`, `strerror_r()`, `strftime()`, `strlen()`, `strncat()`, `strncmp()`, `strncpy()`, `strpbrk()`, `strrchr()`, `strspn()`, `strstr()`, `strtod()`, `strtof()`, `strtoimax()`, `strtok()`, `strtok_r()`, `strtol()`, `strtold()`, `strtoll()`, `strtoul()`, `strtoull()`, `strtoumax()`, `strxfrm()`, `time()`, `tolower()`, `toupper()`, `tzname`, `tzset()`, `va_arg()`, `va_copy()`, `va_end()`, `va_start()`, `vsnprintf()`, `vsprintf()`, `vsscanf()`;

POSIX_SINGLE_PROCESS: `confstr()`, `environ`, `errno`, `getenv()`, `setenv()`, `sysconf()`, `uname()`, `unsetenv()`;

The philosophy of the base RCC profile AEP subset is to allow functions that are simply libraries (rather than OS services), but remove services that could conflict with the lean container execution model in that profile. This basically leaves the typical ANSI-C (or ISO C99) runtime library (without I/O) – most DSPs have this available, even those environments without multithreading.

7 Summary of OpenCPI RCC authoring model

- RCC workers are written to implement the Worker interface, called by the container, and optionally use the Container interface, called by the worker. The container interface is a set of defined container functions.
- RCC workers may call the local services functions defined in the RCC AEP.
- RCC workers use data structures and ordinals as defined in the metadata code generation section (which will be automatically generated from the OWD).
- When not the default, the RCC worker authors must make the following implementation-specific metadata available to the build process (see OWD above) for compiling/linking/loading worker implementation binaries on a given platform:
 - *Implementation name*
 - *Which control operations are implemented*
 - *Which properties require `beforeQuery` or `afterConfigure` notification*
 - *Static memory allocation requirements of the implementation code*
 - *Minimum number of buffers required at each port (default is one)*
- The information in the OWD is used to drive the build process and the code generation process for the RCC worker's header and skeleton files.

7.1 Worker code example

Here is a simple example of a C-language “Xyz” worker whose:

- initial run condition was the default (condition == NULL, usecs == 0, run when all ports are ready, no timeout),
- initialize, release and test methods are empty
- one input port (0) with interface XyzIn (only oneways) and one output port (1) XyzOut (only oneways)
- one oneway IDL interface operation Op1 on input (i.e. can ignore “operation”), which is an array of 100 “shorts”.
- one oneway IDL interface operation Op2 on output
- one simple property, called center_frequency, of type float.

The worker would have automatically generated types and structures like this (based on OCS and OWD), and put in a file called “Xyz_Worker.h”:

```
#include "RCC_Worker.h"
typedef struct { /* structure for defined properties /
    RCCFloat center_frequency;
} XyzProperties;
typedef struct { // structure for message for operation
    int16_t ishorts[100];
} XyzInOp1;

typedef struct { // structure for message for operation
    int16_t oshorts[100];
} XyzOutOp2;

typedef enum { // port ordinals
    XYZ_IN,
    XYZ_OUT
} XyzPort;

typedef enum { // operation ordinals
    XYZ_OUT_OP2
} XyzOutOperation;

typedef enum { // operation ordinals
    XYZ_IN_OP1
} XyzInOperation;
```

The actual code for the worker would look like this:

```

#include "Xyz_Worker.h"

/* Define the initialize method, setting output operation to be a */
/* constant, since it is the only one.*/
static RCCResult
initialize(RCCWorker *w) {
    w->ports[XYZ_OUT].output.u.operation = XYZ_OUT_OP2;
    return RCC_OK;
}

/* Define run method to call the "compute" function, reading from */
/* input buffer, writing to output, applying current value of the */
/* "center frequency" property.*/
static RCCResult
run(RCCWorker *w, RCCBoolean timedout, RCCBoolean *newRunCondition)
{
    XyzProperties *p = w->properties;
    XyzInOp1      *in = w->ports[XYZ_IN].current.data;
    XyzOutOp2     *out = w->ports[XYZ_OUT].current.data;

    /* Do computation based in ishorts, and frequency put results */
    /* in oshorts. Extern is here simply for readability. */
    extern void compute(int16_t *, int16_t *, float);

    compute(in->ishorts, out->oshorts, p->center_frequency);
    /* Ask container to get new input and output buffers */
    return RCC_ADVANCE;
}

/* Initialize dispatch table for container, in a global symbol /
RCCDispatch Xyz = {
    /* Consistency checking attributes */
    RCC_VERSION, 1, 1, sizeof(XyzProperties), RCC_NULL, RCC_FALSE,
    /* Methods */
    initialize, RCC_NULL, RCC_NULL, RCC_NULL, RCC_NULL, RCC_NULL,
    RCC_NULL, run,
    /* Default run condition */
    RCC_NULL
};

```

Other than the compute function, the above example compiles to use less than 120 bytes on a Pentium processor.

A simple non-preemptive single-threaded container implementation would have a loop, testing run conditions, and calling run methods. A more complex environment might run workers in different threads for purposes of time preemption, prioritization, etc. This model allows a variety of container execution models while keeping the worker model simple.

So, on each execution, the worker sees the status of all I/O ports, and can read from current input buffers, and write to current output buffers. It must return to get new buffers, after specifying whether buffers are consumed or filled during the execution.

This simple execution environment can be easily implemented in full-function GPP environments, providing a test environment and a migration path to more minimal embedded environments such as a single-threaded environment with no real operating system.

7.1.1.1 Worker implementation using input callback

This illustrates an input-based callback function that has a sequential coding style in which the worker blocks on the two-way remote call until the response is received and processing can continue. No extra threads are used.

The code is:

```
#include "Xyt_Worker.h"

/* Define the input port callback method to call the local "compute"
 * function, reading from input buffer, applying the value of the
 * "center frequency" property, followed by two way "process"
 * operation, finally writing the result to output port XYT_OUT */
static RCCResult
computeInput(RCCWorker *this, RCCPort *inPort, RCCResult reason)
{
    XytProperties *p = this->properties;
    RCCContainer *c = &this->container;
    XytInOp1 *data = inPort->current.data;
    RCCPort *computeOut = &this->ports[XYT_COMPUTE_REQUEST_OUT],
        *computeIn = &this->ports[XYT_COMPUTE_REPLY_IN],
        *otherOut = &this->ports[XYT_OUT];
    extern void compute(int16_t *, int16_t *, float);

    if (reason != RCC_OK)
        return RCC_FATAL;

    /* Do computation based in is shorts, and frequency; put results
     * back in is shorts, in place */
    compute(data->is shorts, data->is shorts, p->center_frequency);

    /* Call Process operation on "uses" port: buffer ownership passes
     * back to container. */
    c->send(computeOut, &inPort->current, XYT_COMPUTE_PROCESS,
        inPort->input.length);
    /* Wait until response received (or 100 msec error timeout). */
    c->wait(computeIn, 100, 100000);
    if (computeIn->input.u.exception == RCC_NO_EXCEPTION)
        c->send(otherOut, &computeIn->current, XYT_OUT_OP2,
            computeIn->input.length);
    else
        c->release(&computeIn->current);
    return RCC_OK;
}

/* Define the initialize method to perform worker initialization. */
static RCCResult
initialize(RCCWorker *this)
{
    this->ports[XYT_IN].callBack = computeInput;
    return RCC_OK;
}

/* Initialize dispatch table provided to container. We only need the
 * initialize method to register the callback. No run method needed */
RCCDispatch
Xyt = {
    /* Consistency checking attributes */
    RCC_VERSION, 1, 1, sizeof(XytProperties), RCC_NULL, RCC_TRUE,
    /* Methods */
    initialize
    /* all remaining members zero/NULL */
};
```

Other than the compute function, the above example compiles to use less than 300 bytes on a Pentium processor.

7.1.1.2 Worker implementation using the using state-machine style

This illustrates a finite state machine coding style in a worker that maintains an internal state to simulate blocking on a two-way remote call until the response is received and processing can continue.

```

#include "Xyt_Worker.h"
/* Define two different run conditions to represent two states */
static uint32_t
    state1Ports[] = {1 << XYT_IN, 0},
    state2Ports[] = {1 << XYT_COMPUTE_REPLY_IN, 0};
static RCCRunCondition
    awaitingInput = {state1Ports},
    awaitingResponse = {state2Ports};

/* Define start method, set run condition, which is also state. /
static RCCResult
initialize(RCCWorker this)
{
    this->runCondition = &awaitingInput;
    return RCC_OK;
}
/* Define run method to call the local "compute" function
 * reading from input buffer, applying current value of the
 * "center frequency" property, followed by remote two way
 * "process" operation, finally writing result to output. */
static RCCResult
run(RCCWorker *this, RCCBoolean timedout, RCCBoolean *newRunCondition)
{
    RCCContainer *c = &this->container;

    /* Use run condition as state indicator */
    if (this->runCondition == &awaitingInput) {
        RCCPort
            *inPort      = &this->ports[XYT_IN],
            *computeOut   = &this->ports[XYT_COMPUTE_REQUEST_OUT];
        XytInOp1 *in      = inPort->current.data;
        XytProperties *p = this->properties;

        /* do some computation based in ishorts, and frequency;
         * put results back in same buffer (in-place) */
        extern void compute(int16_t *, int16_t *, float);
        compute(in->ishorts, in->ishorts, p->center_frequency);

        /* Call Process op on user port - buffer ownership passes back
         * to container. Input port is advanced by taking buffer away
         * from it. */
        c->send(computeOut, &inPort->current, XYT_COMPUTE_PROCESS,
            inPort->input.length); /* length of message */
        this->runCondition = &awaitingResponse; /* update state */
    } else {
        RCCPort *computeIn = &this->ports[XYT_COMPUTE_REPLY_IN];
        RCCPort *otherOut = &this->ports[XYT_OUT];

        if (computeIn->input.u.exception == 0)
            c->send(otherOut, &computeIn->current, XYT_OUT_OP2,
                computeIn->input.length);
        else
            c->advance(computeIn, 0);
        this->runCondition = &awaitingInput; /* update state */
    }
    *newRunCondition = RCC_TRUE; /* to container: new runcondition /
    return RCC_OK;
}
/* continued on next page */

```

```

/* Initialize dispatch table provided to container. We only need the
 * initialize method to register the callback. No run method needed */
RCCDispatch Xyt = {
    /* Consistency checking attributes */
    RCC_VERSION, 1, 1, sizeof(XytProperties), RCC_NULL, RCC_FALSE,
    /* Methods */
    initialize, RCC_NULL, RCC_NULL, RCC_NULL, RCC_NULL, RCC_NULL,
    RCC_NULL, run,
    /* all remaining members zero/NULL */
};

```

Other than the compute function, the above example compiles to use less than 330 bytes on a Pentium processor.

8 Glossary

Configuration Properties – Named values associated with a worker that may be read or written by the control application and/or the worker. Their values indicate or control aspects of the worker’s operation. Reading and writing these property values may or may not have side effects on the operation of the worker. Configuration properties with side effects can be used for custom worker control. Each worker may have its own, possibly unique, set of configuration properties. Some properties come from the OCS: they are common to all implementations that implement that OCS. Other properties can be added in the OWD that are specific to that particular worker implementation.

Container – An OpenCPI infrastructure element that “contains”, manages and executes a set of application workers. Logically, the container “surrounds” the workers, mediating all interactions between the worker and the rest of the system.

Worker Attribute – An attribute related to a particular implementation (design) of a worker, for one of its interfaces. I.e. one that is not necessarily common across a set of implementations of the same high level component definition (OCS).

Worker – A concrete implementation (and possibly runtime instance) of a component, generally existing within a container. A worker is implemented consistent with its OWD.

9 References

ID	Link/URL
1	OpenCPI Component Development Guide: https://github.com/opencpi/opencpi/raw/master/doc/pdf/OpenCPI_Component_Development.pdf

10 List of Abbreviations and Acronyms

ACI	Application Control Interface
ADC	Analog-to-Digital Converter
API	Application Programming Interface
CBD	Component-Based Development
CDK	Component Developer's Kit
CORBA	Common Object Request Broker Architecture
COTS	Commercial Off The Shelf
DAC	Digital to Analog Converter
DDS	Data Distribution Service
DMA	Direct Memory Access
DoD	Department of Defense
DSP	Digital Signal Processor
FGM	Fine-Grained Multicore
FPGA	Field Programmable Gate Array
GBE	GigaBit Ethernet
GPGPU	General Purpose computing on a Graphics Processing Unit
GPP	General Purpose Processor
GPU	Graphics Processing Unit
HCI	Human/Computer Interface
HDL	Hardware Description Language
OpenCL	Open Computing Language
OpenCPI	Open-Source Component Portability Infrastructure
OpenCV	Open Computer Vision library
RDMA	Remote Direct Memory Access
SCA	Software Communications Architecture
SOW	Statement of Work

SWaP	Size, Weight and Power
VHDL	VHSIC Hardware Description Language
XML	Extensible Markup Language